

Final Year B. Tech, Sem VII 2022-23

PRN – 2020BTECS00211

Name – Aashita Narendra Gupta

High Performance Computing

Lab Batch: B4

SLIP No - 42

HPC LAB ESE

Github Link for Code - https://github.com/Aashita06/HPC_Practicals

Q.1) Write a program to demonstrate distributed sum of an array using MPI.

→

Code:

```
#include "mpi.h"
#include <stdio.h>

#define localSize 1000

// store the subarray data coming from process 0;
int local[1000];

int main(int argc, char **argv)
{
    int rank;
    int num;

    int n = 20;
    int arr[20] = {1, 2, 3, 4, 5, 6, 7, 8, 9,
10,11,12,13,14,15,16,17,18,19,20};

    int per_process, elements_received;

    MPI_Init(&argc, &argv);

    MPI_Comm_size(MPI_COMM_WORLD, &num);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    MPI_Status status;
```

```

    // process with rank 0 will divide data among all processes and add
    partial sums to get final sum
    if (rank == 0)
    {
        int index, i;

        per_process = n / num;

        double start = MPI_Wtime();

        // if more than 1 processes available
        if (num > 1)
        {
            //divide array data among processes
            for (i = 1; i < num - 1; i++)
            {
                //calculating first index of subarray that need to be send to
                ith process
                index = i * per_process;

                //send no of elements and subarray of that lenght to each
                process
                MPI_Send(&per_process, 1, MPI_INT, i, 0, MPI_COMM_WORLD);
                MPI_Send(&arr[index], per_process, MPI_INT, i, 0,
                MPI_COMM_WORLD);
            }

            // for last process send all remaining elements
            index = i * per_process;
            int ele_left = n - index;

            MPI_Send(&ele_left, 1, MPI_INT, i, 0, MPI_COMM_WORLD);
            MPI_Send(&arr[index], ele_left, MPI_INT, i, 0, MPI_COMM_WORLD);
        }

        // add numbers on process with rank 0
        int sum = 0;
        for (int i = 0; i < per_process; i++)
        {
            sum += arr[i];
        }

        // add all partial sums from all processes
        int tmp;
        for (int i = 1; i < num; i++)
        {
            MPI_Recv(&tmp, 1, MPI_INT, MPI_ANY_SOURCE, 0, MPI_COMM_WORLD,
            &status);

```

```

        int sender = status.MPI_SOURCE;

        sum += tmp;
    }

    printf("\nSum of array = %d\n", sum);

    double end = MPI_Wtime();
    printf("Time required by %d processors : %f", num, end-start);
}
else // if rank of process is not 0, then receive elements and calculate
partial sums
{
    // receive no of elements and elements form process 0 and store them
on local array
    MPI_Recv(&elements_received, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
&status);

    MPI_Recv(&local, elements_received, MPI_INT, 0, 0, MPI_COMM_WORLD,
&status);

    // calculate partial local sum
    int partial_sum = 0;
    for (int i = 0; i < elements_received; i++)
    {
        partial_sum += local[i];
        // printf("\nPartial Sum of rank %d is %d ", rank, partial_sum);
    }

    //send calculated partial sum to process with rank 0
    MPI_Send(&partial_sum, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
}

MPI_Finalize();
return 0;
}


```

Output:

```
PROBLEMS  OUTPUT  TERMINAL  GITLENS  DEBUG CONSOLE
PS C:\Users\Ashitra\OneDrive\Desktop\Lab_Prog> mpiexec -n 2 ./sumarray.exe
Sum of array = 210
Time required by 2 processors : 0.001492
PS C:\Users\Ashitra\OneDrive\Desktop\Lab_Prog> mpiexec -n 4 ./sumarray.exe
Sum of array = 210
Time required by 4 processors : 0.001097
PS C:\Users\Ashitra\OneDrive\Desktop\Lab_Prog> mpiexec -n 6 ./sumarray.exe
Sum of array = 210
Time required by 6 processors : 0.001488
PS C:\Users\Ashitra\OneDrive\Desktop\Lab_Prog> mpiexec -n 8 ./sumarray.exe
Sum of array = 210
Time required by 8 processors : 0.002153
PS C:\Users\Ashitra\OneDrive\Desktop\Lab_Prog> mpiexec -n 10 ./sumarray.exe
Sum of array = 210
Time required by 10 processors : 0.003042
PS C:\Users\Ashitra\OneDrive\Desktop\Lab_Prog> 
```

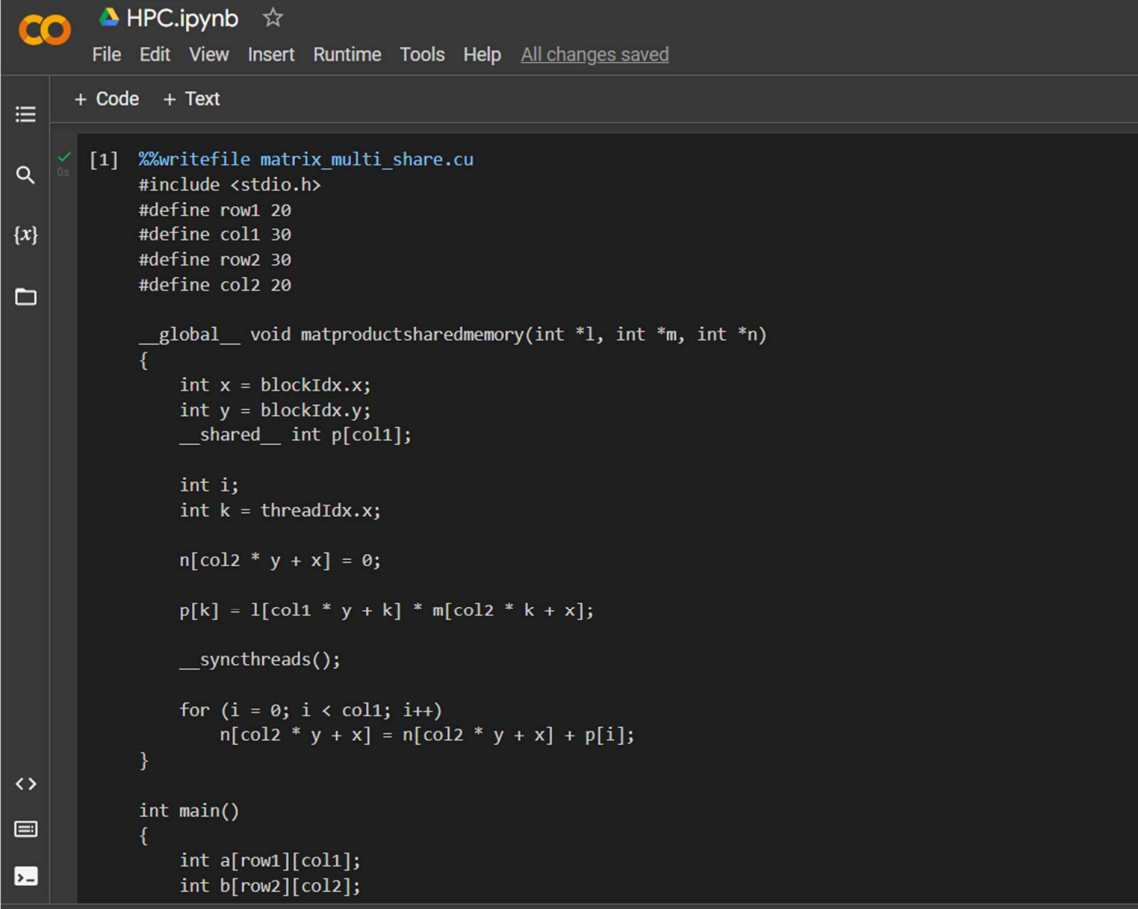
We are analysis the execution time taken by different number of processors to calculate the distributed sum of array. Here, due to communication overhead, time is increasing with increasing processors.

Processors	Time
2	0.001492
4	0.001097
6	0.001488
8	0.002153
10	0.003042



Q.2) Implement matrix matrix multiplication using CUDA (shared memory).
(Conduct performance and speedup analysis of all programs)
→

Code:



The screenshot shows a Jupyter Notebook interface with a dark theme. The top bar includes the Jupyter logo, the text 'HPC.ipynb', and a star icon. Below this is a menu bar with 'File', 'Edit', 'View', 'Insert', 'Runtime', 'Tools', and 'Help', followed by the status 'All changes saved'. The left sidebar contains icons for a menu, search, variables, file explorer, and input/output. The main area displays a code cell with the following CUDA code:

```
[1] %%writefile matrix_multi_share.cu
#include <stdio.h>
#define row1 20
#define col1 30
#define row2 30
#define col2 20

__global__ void matproductsharedmemory(int *l, int *m, int *n)
{
    int x = blockIdx.x;
    int y = blockIdx.y;
    __shared__ int p[col1];

    int i;
    int k = threadIdx.x;

    n[col2 * y + x] = 0;

    p[k] = l[col1 * y + k] * m[col2 * k + x];

    __syncthreads();

    for (i = 0; i < col1; i++)
        n[col2 * y + x] = n[col2 * y + x] + p[i];
}

int main()
{
    int a[row1][col1];
    int b[row2][col2];
```

```
+ Code + Text
[1] int c[row1][col2];
    int *d, *e, *f;
    int i, j;

    for (i = 0; i < row1; i++)
    {
        for (j = 0; j < col1; j++)
        {
            a[i][j] = 2;
        }
    }

    for (i = 0; i < row2; i++)
    {
        for (j = 0; j < col2; j++)
        {
            b[i][j] = 3;
        }
    }

    cudaMalloc((void **)&d, row1 * col1 * sizeof(int));
    cudaMalloc((void **)&e, row2 * col2 * sizeof(int));
    cudaMalloc((void **)&f, row1 * col2 * sizeof(int));

    cudaMemcpy(d, a, row1 * col1 * sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(e, b, row2 * col2 * sizeof(int), cudaMemcpyHostToDevice);

    dim3 grid(col2, row1);

    matproductsharedmemory<<<grid, col1>>>(d, e, f);
    cudaDeviceSynchronize();
```

```
+ Code + Text
[1] cudaMemcpy(c, f, row1 * col2 * sizeof(int), cudaMemcpyDeviceToHost);

    for (i = 0; i < row1; i++)
    {
        for (j = 0; j < col2; j++)
        {
            if (c[i][j] != 180)
            {
                printf("False\n");
                return -1;
            }
        }
    }

    cudaFree(d);
    cudaFree(e);
    cudaFree(f);

    printf("True\n");
    return 0;
}
```

Output:

```
Writing matrix_multi_share.cu

Invcc -o matrix_multi_share matrix_multi_share.cu
!./matrix_multi_share

True
```

Profiling of matrix matrix multiplication program in CUDA (Shared Memory).

```
Insys profile --stats=true matrixshare

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.
Collecting data...
True
The target application terminated with signal 11 (SIGSEGV)
Processing events...
Capturing symbol files...
Saving temporary "/tmp/nsys-report-b443-e0d0-593d-c942.qdstrm" file to disk...
Creating final output files...

Processing [=====100%]
Saved report file to "/tmp/nsys-report-b443-e0d0-593d-c942.qdrep"
Exporting 11741 events: [=====100%]

Exported successfully to
/tmp/nsys-report-b443-e0d0-593d-c942.sqlite

CUDA API Statistics:

Time(%) Total Time (ns) Num Calls Average Minimum Maximum Name
-----
91.4 343,509,508 3 114,503,169.3 172,494 343,158,687 cudaMalloc
7.5 28,200,933 3 9,400,311.0 8,871,599 9,717,737 cudaMemcpy
0.9 3,492,907 3 1,164,302.3 229,326 2,088,871 cudaFree
0.1 497,851 2 248,925.5 4,339 493,512 cudaDeviceSynchronize
0.0 48,551 1 48,551.0 48,551 48,551 cudaLaunchKernel

CUDA Kernel Statistics:
```

```
+ Code + Text

CUDA Kernel Statistics:

Time(%) Total Time (ns) Instances Average Minimum Maximum Name
-----
100.0 478,557 1 478,557.0 478,557 478,557 kernel_sum(float const*, float const*, float*, int)

CUDA Memory Operation Statistics (by time):

Time(%) Total Time (ns) Operations Average Minimum Maximum Operation
-----
68.6 18,797,709 2 9,398,854.5 9,373,191 9,424,518 [CUDA memcpy HtoD]
31.4 8,603,755 1 8,603,755.0 8,603,755 8,603,755 [CUDA memcpy DtoH]

CUDA Memory Operation Statistics (by size in KiB):

Total Operations Average Minimum Maximum Operation
-----
78,125.000 2 39,062.500 39,062.500 39,062.500 [CUDA memcpy HtoD]
39,062.500 1 39,062.500 39,062.500 39,062.500 [CUDA memcpy DtoH]
```

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
91.4	343,509,508	3	114,503,169.3	172,494	343,158,687	cudaMalloc
7.5	28,200,933	3	9,400,311.0	8,871,599	9,717,737	cudaMemcpy
0.9	3,492,907	3	1,164,302.3	229,326	2,088,871	cudaFree
0.1	497,851	2	248,925.5	4,339	493,512	cudaDeviceSynchronize
0.0	48,551	1	48,551.0	48,551	48,551	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	478,557	1	478,557.0	478,557	478,557	kernel_sum(float const*, float const*, float*, int)

CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
68.6	18,797,709	2	9,398,854.5	9,373,191	9,424,518	[CUDA memcpy HtoD]
31.4	8,603,755	1	8,603,755.0	8,603,755	8,603,755	[CUDA memcpy DtoH]

CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
78,125.000	2	39,062.500	39,062.500	39,062.500	[CUDA memcpy HtoD]
39,062.500	1	39,062.500	39,062.500	39,062.500	[CUDA memcpy DtoH]

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
64.4	401,333,339	15	26,755,555.9	55,150	100,203,011	poll
35.2	219,484,441	674	325,644.6	1,055	117,154,957	ioctl
0.2	991,103	82	12,086.6	3,099	33,028	open64
0.1	531,816	10	53,181.6	19,174	75,883	sem_timedwait
0.0	223,782	4	55,945.5	40,122	64,325	pthread_create
0.0	151,659	25	6,066.4	1,673	20,746	fopen
0.0	107,313	3	35,771.0	31,945	42,784	fgets
0.0	74,895	11	6,808.6	3,708	9,982	write
0.0	65,372	18	3,631.8	1,003	36,337	fclose
0.0	63,336	6	10,556.0	3,432	18,294	open
0.0	36,494	6	6,082.3	1,868	10,893	fgetc
0.0	35,079	1	35,079.0	35,079	35,079	sem_wait
0.0	16,371	10	1,637.1	1,029	2,269	read
0.0	14,790	2	7,395.0	6,809	7,981	socket
0.0	11,699	4	2,924.8	1,151	4,835	fcntl
0.0	9,580	1	9,580.0	9,580	9,580	pipe2
0.0	8,326	1	8,326.0	8,326	8,326	connect
0.0	7,228	2	3,614.0	3,118	4,110	fread
0.0	2,295	2	1,147.5	1,147	1,148	msync
0.0	1,531	1	1,531.0	1,531	1,531	bind
0.0	1,193	1	1,193.0	1,193	1,193	listen

Report file moved to "/content/report3.qdrep"
Report file moved to "/content/report3.sqlite"

Keeping No. of blocks constants and increasing no. of threads, we are calculating speedup of the following pmatrix multiplication program.

Speedup	No.of threads
3.6, 1.04, 1.0, 1.31	8
3.7, 1.32, 1.2, 1.51	16
4.23, 2.5, 2.1, 2.7	32

